

Making asynchronous system calls in Linux kernel

Student's Individual Study
Sakhybekova Balnur

"Asynchronous I/O" allows for a process to monitor multiple file descriptors and get notifications when I/O is possible on them.

A level-triggered AIO mechanism provides (when queried) information about which AIO events are still pending. In general, this translates to a set of file descriptors on which reading (or writing) can be performed without blocking.

An edge-triggered mechanism on the other hand provides information about which events have changed status (from non-pending to pending) since the last query.

Dealing with asynchronous events on Linux

- • Threading
- Signals
- The SIGIO signal
- `select()` and `poll()`
- `epoll()`
- POSIX asynchronous I/O (AIO)

epoll is a Linux kernel system call for a scalable I/O event notification mechanism, first introduced in version 2.5.44 of the Linux kernel mainline.^[1] Its function is to monitor multiple file descriptors to see whether I/O is possible on any of them. It is meant to replace the older POSIX **select()** and **poll()** system calls, to achieve better performance in more demanding applications, where the number of watched file descriptors is large (unlike the older system calls, which operate in $O(n)$ time, **epoll** operates in $O(1)$ time^[2]).

Makefile

```
all: epoll_example
```

```
epoll_example: epoll_example.c  
gcc -Wall -Werror -o $@  
epoll_example.c
```

```
clean:
```

```
@rm -v epoll_example
```

Used libraries

```
#include <stdio.h> // for fprintf()
```

```
#include <unistd.h> // for close(), read()
```

```
#include <sys/epoll.h> // for epoll_create1(), epoll_ctl(), struct epoll_event
```

```
#include <string.h> // for strncmp
```

```
1. #include <stdio.h>      // for fprintf()
2. #include <unistd.h>     // for close()
3. #include <sys/epoll.h>   // for epoll_create1()
4.
5. int main()
6. {
7.     int epoll_fd = epoll_create1(0);
8.
9.     if(epoll_fd == -1)
10.    {
11.        fprintf(stderr, "Failed to create epoll file descriptor\n");
12.        return 1;
13.    }
14.
15.    if(close(epoll_fd) )
16.    {
17.        fprintf(stderr, "Failed to close epoll file descriptor\n");
18.        return 1;
19.    }
20.    return 0;
21. }
```

```
1. #include <stdio.h>      // for fprintf()
2. #include <unistd.h>     // for close()
3. #include <sys/epoll.h>   // for epoll_create1(), epoll_ctl(), struct
   epoll_event
4.
5. int main()
6. {
7.     struct epoll_event event;
8.     int epoll_fd = epoll_create1(0);
9.
10.    if(epoll_fd == -1)
11.    {
12.        fprintf(stderr, "Failed to create epoll file descriptor\n");
13.        return 1;
14.    }
15.
16.    event.events = EPOLLIN;
17.    event.data.fd = 0;
18.
19.    if(epoll_ctl(epoll_fd, EPOLL_CTL_ADD, 0, &event))
20.    {
21.        fprintf(stderr, "Failed to add file descriptor to epoll\n");
22.        close(epoll_fd);
23.        return 1;
24.    }
25.
26.    if(close(epoll_fd))
27.    {
28.        fprintf(stderr, "Failed to close epoll file descriptor\n");
29.        return 1;
30.    }
31.    return 0;
32. }
```

```
8. int main()
9. {
10.     int running = 1, event_count, i;
11.     size_t bytes_read;
12.     char read_buffer[READ_SIZE + 1];
13.     struct epoll_event event, events[MAX_EVENTS];
14.     int epoll_fd = epoll_create1(0);
15.
16.     if(epoll_fd == -1)
17.     {
18.         fprintf(stderr, "Failed to create epoll file descriptor\n");
19.         return 1;
20.     }
21.
22.     event.events = EPOLLIN;
23.     event.data.fd = 0;
24.
25.     if(epoll_ctl(epoll_fd, EPOLL_CTL_ADD, 0, &event))
26.     {
27.         fprintf(stderr, "Failed to add file descriptor to epoll\n");
28.         close(epoll_fd);
29.         return 1;
30.     }
31.
32.     while(running)
33.     {
34.         printf("\nPolling for input...\n");
35.         event_count = epoll_wait(epoll_fd, events, MAX_EVENTS, 30000);
36.         printf("%d ready events\n", event_count);
37.         for(i = 0; i < event_count; i++)
38.         {
39.             printf("Reading file descriptor '%d' -- ", events[i].data.fd);
40.             bytes_read = read(events[i].data.fd, read_buffer, READ_SIZE);
41.             printf("%zd bytes read.\n", bytes_read);
42.             read_buffer[bytes_read] = '\0';
43.             printf("Read '%s'\n", read_buffer);
44.
45.             if(!strncmp(read_buffer, "stop\n", 5))
46.                 running = 0;
47.         }
48.     }
49.
50.     if(close(epoll_fd))
51.     {
52.         fprintf(stderr, "Failed to close epoll file descriptor\n");
53.         return 1;
54.     }
55.     return 0;
56. }
```

sysprog [Running]

Cc 18:42

Activities Terminal

bestriess@bestriess-VirtualBox: ~/sysprog/sis

```
File Edit View Search Terminal Help
epoll_example.c Makefile
bestriess@bestriess-VirtualBox:~/sysprog/sis$ ./epoll_example
bash: ./epoll_example: No such file or directory
bestriess@bestriess-VirtualBox:~/sysprog/sis$ make
Makefile:4: *** missing separator. Stop.
bestriess@bestriess-VirtualBox:~/sysprog/sis$ ls
epoll_example.c Makefile
bestriess@bestriess-VirtualBox:~/sysprog/sis$ gedit Makefile
bestriess@bestriess-VirtualBox:~/sysprog/sis$ make
Makefile:4: *** missing separator. Stop.
bestriess@bestriess-VirtualBox:~/sysprog/sis$ gcc -o epoll_example epoll_example.c
gcc: fatal error: no input files
compilation terminated.
bestriess@bestriess-VirtualBox:~/sysprog/sis$ gcc -o epoll_example epoll_example.c
bestriess@bestriess-VirtualBox:~/sysprog/sis$ ls
epoll_example epoll_example.c Makefile
bestriess@bestriess-VirtualBox:~/sysprog/sis$ ./epoll_example

Polling for input...
0 ready events

Polling for input...
```

Activities Terminal Cc 18:42

bestriess@bestriess-VirtualBox: ~/sysprog/sis

```
File Edit View Search Terminal Help
le.c
bestriess@bestriess-VirtualBox:~/sysprog/sis$ ls
epoll_example epoll_example.c Makefile
bestriess@bestriess-VirtualBox:~/sysprog/sis$ ./epoll_example

Polling for input...
0 ready events

Polling for input...

1 ready events
Reading file descriptor '0' -- 1 bytes read.
Read ' '
'

Polling for input...
hello
1 ready events
Reading file descriptor '0' -- 6 bytes read.
Read 'hello'
'

Polling for input...
'
```

Left ☰

sysprog [Running]
Activities Terminal Cc 18:43

```
bestriess@bestriess-VirtualBox: ~/sysprog/sis
File Edit View Search Terminal Help
Polling for input...
By Balnur Sakhybekovaaa
1 ready events
Reading file descriptor '0' -- 10 bytes read.
Read 'By Balnur '

Polling for input...
1 ready events
Reading file descriptor '0' -- 10 bytes read.
Read 'Sakhybekov'

Polling for input...
1 ready events
Reading file descriptor '0' -- 4 bytes read.
Read 'aaa'
'

Polling for input...
stop
1 ready events
Reading file descriptor '0' -- 5 bytes read.
Read 'stop'
'

bestriess@bestriess-VirtualBox:~/sysprog/sis$
```